



EclipseWorld 2007

November 6-8, 2007

Eclipse Persistence Services The Full Monty

Peter Krogh, Oracle Canada

EclipseLink Project co-Lead

Doug Clarke, Oracle Canada

EclipseLink Project co-Lead

What you will learn

- What the Eclipse Persistence Services Project is
- How this project can be used and its benefits
- Why you will want to use this project
- Usages with Spring
- How you can get involved

Eclipse Persistence Services

- Eclipse runtime project
 - Nicknamed “EclipseLink”
 - Currently Incubating in Technology Project
- Comprehensive
 - EclipseLink JPA: Object-Relational
 - EclipseLink MOXy: Object-XML
 - EclipseLink SDO: Service Data Objects
 - EclipseLink DBWS: Database Web Services
 - EclipseLink EIS: Non-Relational using JCA
- Defining blueprints for OSGi persistence services

What is Eclipse?

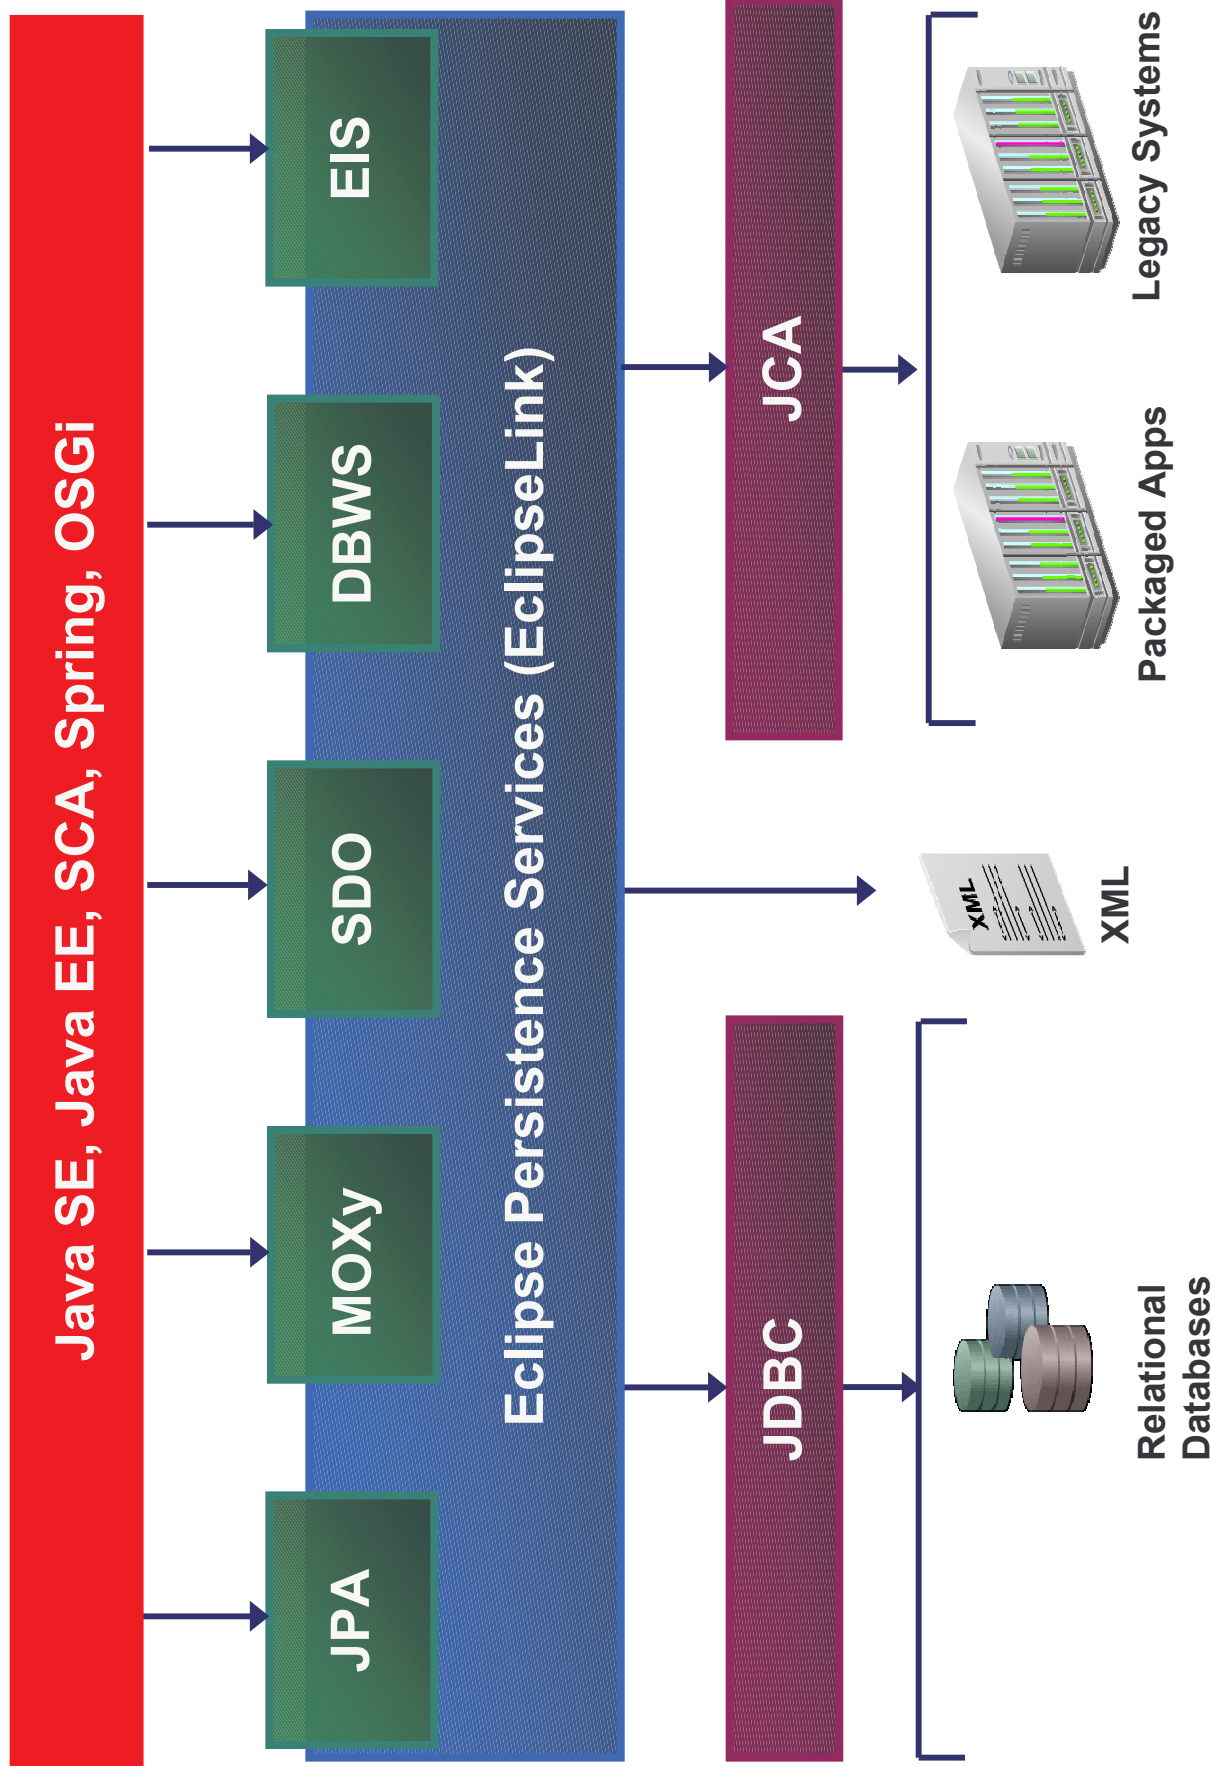
- Eclipse is an open source community
- Eclipse is more than just an IDE
 - Equinox (OSGi), Rich Client Platform (RCP), Higgins (Trust Framework), ...
 - Incubating
 - Maya (Deployment Framework)
 - Swordfish (SOA)
 - **Persistence Services Project (EclipseLink)**
 - Proposals
 - Rich Server Platform, ...

Why Eclipse?

- Eclipse has a strong and vibrant community with an effective governance model
- Good reputation for quality
- Interest from within the Eclipse ecosystem
- Oracle has had a positive experience with its existing participation in Eclipse projects
 - Projects lead by Oracle: Dali, BPEL, JSF
 - Other Oracle contributions: WTP and DTP

Oracle TopLink





Importance

- First comprehensive open source persistence solution
 - Object-Relational and much more
- Based upon product with 12 years of commercial usage
- Shared infrastructure
 - Easily share the same domain model with multiple persistence technologies
 - Leverage metadata for multiple services
- Important part of the Eclipse Ecosystem

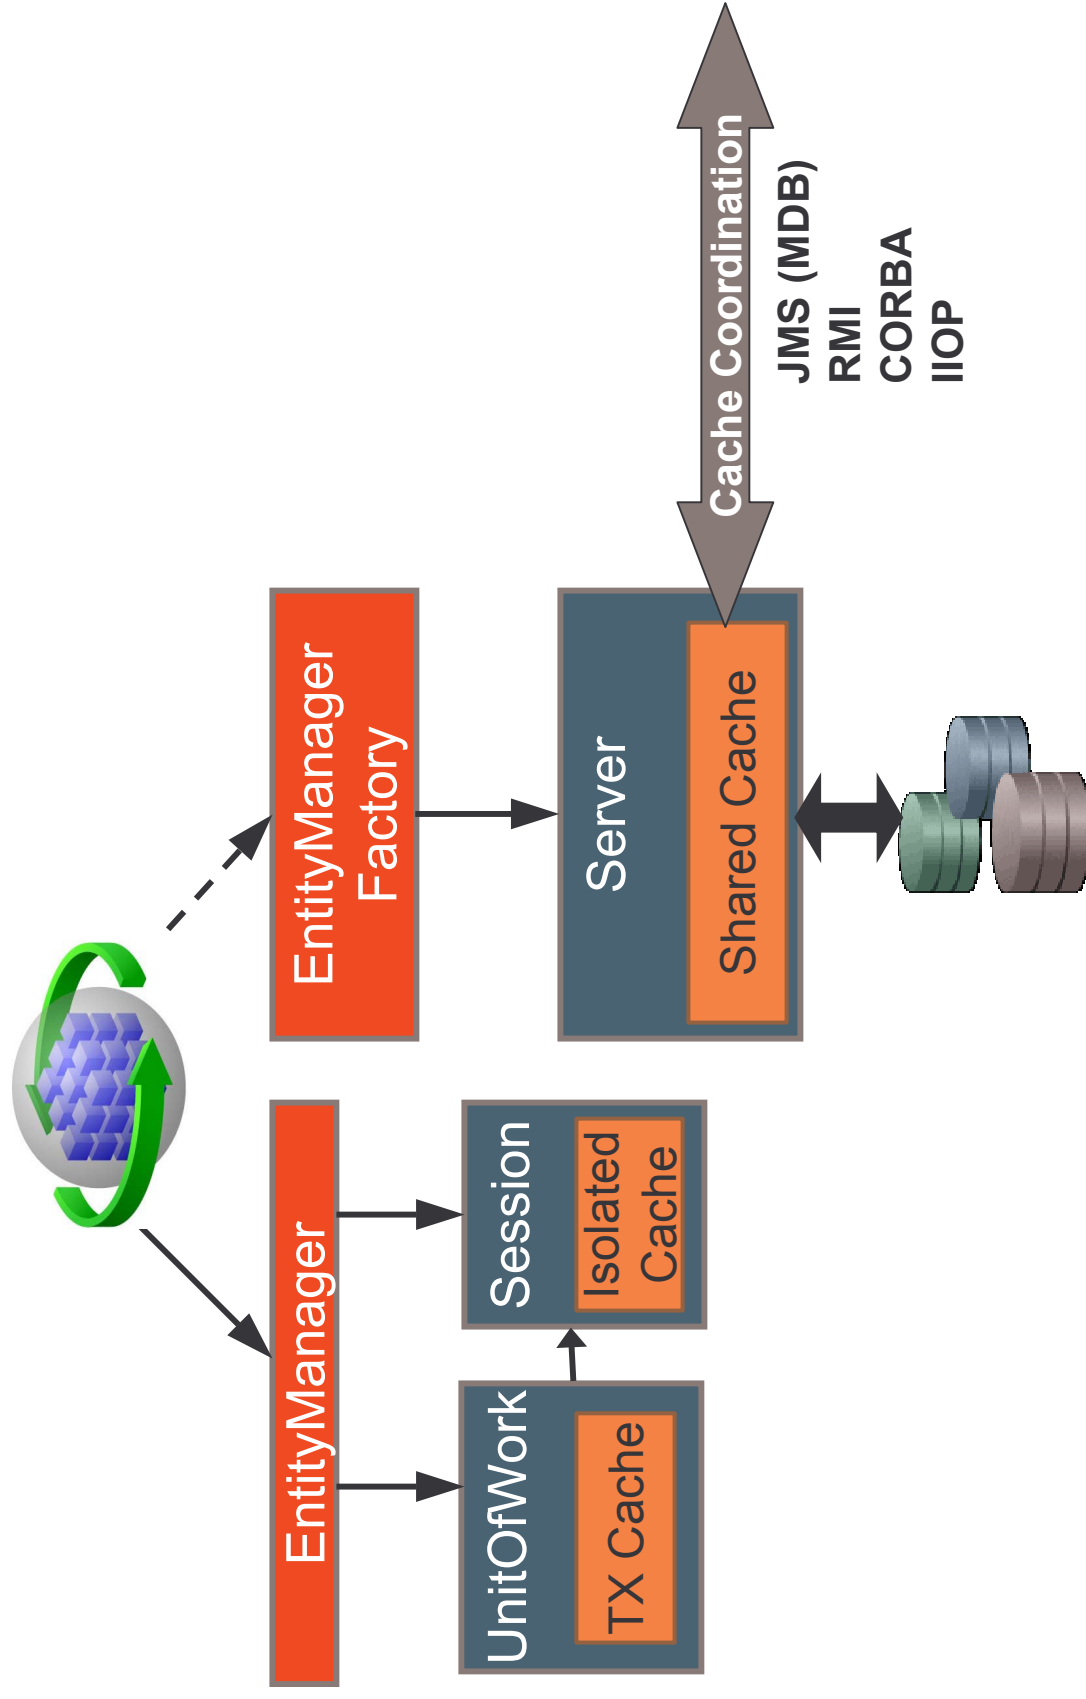
EclipseLink JPA

- JPA 1.0 compliant implementation
- Java EE, Java SE, Web, Spring, and OSGi
- Any JDBC/SQL compliant database
- Extensible and pluggable
- Schema generation
- Key infrastructure:
 - Caching, Locking, Query Framework, Mapping, ...
- ... plus many valuable advanced features

EclipseLink Caching

- Entity caching
 - L2 shared across transactions/users
 - Coordination in a clustered deployment
- Application specific configuration
 - Cache isolation: per client (EM) or shared
 - Cache Type and Size: Weak, Soft-Weak, Full, None
 - Expiration/Invalidation
 - Time to live, Time of day, API
 - Coordination (cluster-messaging)
 - Messaging: JMS, RMI, CORBA, RMI-IIOP, ...
 - Mode: SYNC, SYNC+NEW, INVALIDATE, NONE

Caching Architecture



Configuring the Cache

- Default: objects read are cached and trusted
- Configuration by entity type important
 - Volatility of data
 - Shared usage of data
- Configuration Parameters
 - Cache isolation, type, size, expiry, coordination
 - Refreshing
 - By query (use-case) or descriptor (always)
- Locking is the only way to avoid potential data corruption in concurrent write scenarios

Locking

- Prevent data corruption !!!
- Java Developers think of locking at the object level
- Databases may need to manage locking across many applications
- EclipseLink is able to respect and participate in locks at database level
 - Optimistic: Numeric, Timestamp, All fields, Selected fields, Changed field
 - Pessimistic

Query Framework

- Queries can be defined using
 - Entity Model: JPQL, Expressions, Query-by-example
 - Database: SQL, Stored Procedures
- Customizable
 - Locking, Cache Usage, Refreshing
 - Optimizations: Joining, Batching, parameter binding
 - Result shaping/conversions
- Static or Dynamic
 - Stored Procedure support

EclipseLink JPA Extensions

- Extensions supported through annotations and XML
- Mapping
 - @BasicMap, @BasicCollection, @PrivateOwned, @JoinFetch
 - @Converter, @TypeConverter, @ObjectTypeConverter
- @Cache
 - type, size, isolated, expiry, refresh, cache usage, coordination
 - Cache usage and refresh query hints
- @NamedStoredProcedureQuery
 - IN/OUT/INOUT parameters, multiple cursor results

EclipseLink JPA Extensions

- Locking
 - Non-intrusive policies `@OptimisticLocking`
 - Pessimistic query hints
- JDBC Connection Pooling
- Logging: Diagnostics, SQL, Debugging
- Weaving for lazy fetch and change tracking
 - Dynamic and Static
- Customization
 - Entity Descriptor: `@Customizer`, `@ReadOnly`
 - Session Customizer

Mapping Extensions

```
@Entity
@Cache(type=SOFT_WEAK, coordinationType=SEND_OBJECT_CHANGES)
@OptimisticLocking(type=CHANGED_COLUMNS)
@Converter(name="money", converterClass=MoneyConverter.class)
public class Employee {
    @Id
    private int id;

    private String name;

    @OneToMany(mappedBy="owner")
    @PrivateOwned
    private List<PhoneNumbers> phones;

    @Convert("money")
    private Money salary

    ...
}
```

Database Platform

- Native SQL (dialect) support with custom operators
- Stored Procedure & Function
- Extensible Advanced Data Types support (Struct)
- Database Security
 - Oracle DB's VPD/OLS and Proxy Authentication
- Configurable value return from write
- Supported platforms (default = Auto)
 - MySQL, Derby, Oracle, DB2, Sybase, SQLServer, TimesTen, PostgreSQL, SQLAnywhere, HSQL, Informix, ...

Server Platform

- Simplified configuration and mediator for host container environment
- Enables
 - Direct JTA integration
 - Data Source/JDBC connection un-wrapping
 - JMX MBean deployment
 - Logging integration
- Current Server Platforms
 - SunAS/GlassFish, OracleAS/OC4J, WLS, WAS, JBoss

Performance and Tuning


- Highly configurable and tunable
 - Principle: minimize and optimize database calls
 - Enable application specific tuning
- Flexibility allows efficient business models and relational schemas to be used
- Leverages underlying performance tuning features
 - Java, JDBC and the underlying database technology

EclipseLink JPA Config

- JPA (portable)
 - Persistence.xml with EclipseLink properties
 - Mapping: Annotations and/or orm.xml
 - Query hints
- EclipseLink
 - Sessions Configuration (sessions.xml)
 - Mapping using XML or Code
- EclipseLink JPA
 - JPA + EclipseLink configurations options
 - EclipseLink annotations

EclipseLink MOxy

- Provides complete Object-XML mapping
 - Allows developers to work with XML as objects
 - Efficiently produce and consume XML
 - Document Preservation
- Supports Object-XML standard - JAXB
 - Provides additional flexibility to allow complete control on how objects are mapped



EclipseLink MOXy Benefits

- Rich set of mappings providing complete control and flexibility to map objects to any XSD
 - Direct, composite object, composite collection, inheritance, positional, path, transformation
- Development Approaches
 - Model + Annotations → XSD
 - XSD → Model + Annotations
 - Model + Mappings(Annotations or XML)
- Supports any JAXP compliant parser
 - SAX, DOM, StAX
- Visual Mapping support using Workbench

EclipseLink MOxy: JAXB

```
JAXBContext ctx = JAXBContext.newInstance(classes) ;  
Marshaller marshaller = ctx.createMarshaller() ;  
  
Customer customer = new Customer() ;  
customer.setFirstName("William") ;  
customer.setLastName("Gibson") ;  
  
marshaller.marshal(customer, System.out) ;
```

jaxb.properties:

```
javax.xml.bind.context.factory =  
org.eclipse.persistence.jaxb.JAXBContextFactory
```




EclipseWorld 2007
November 6-8, 2007

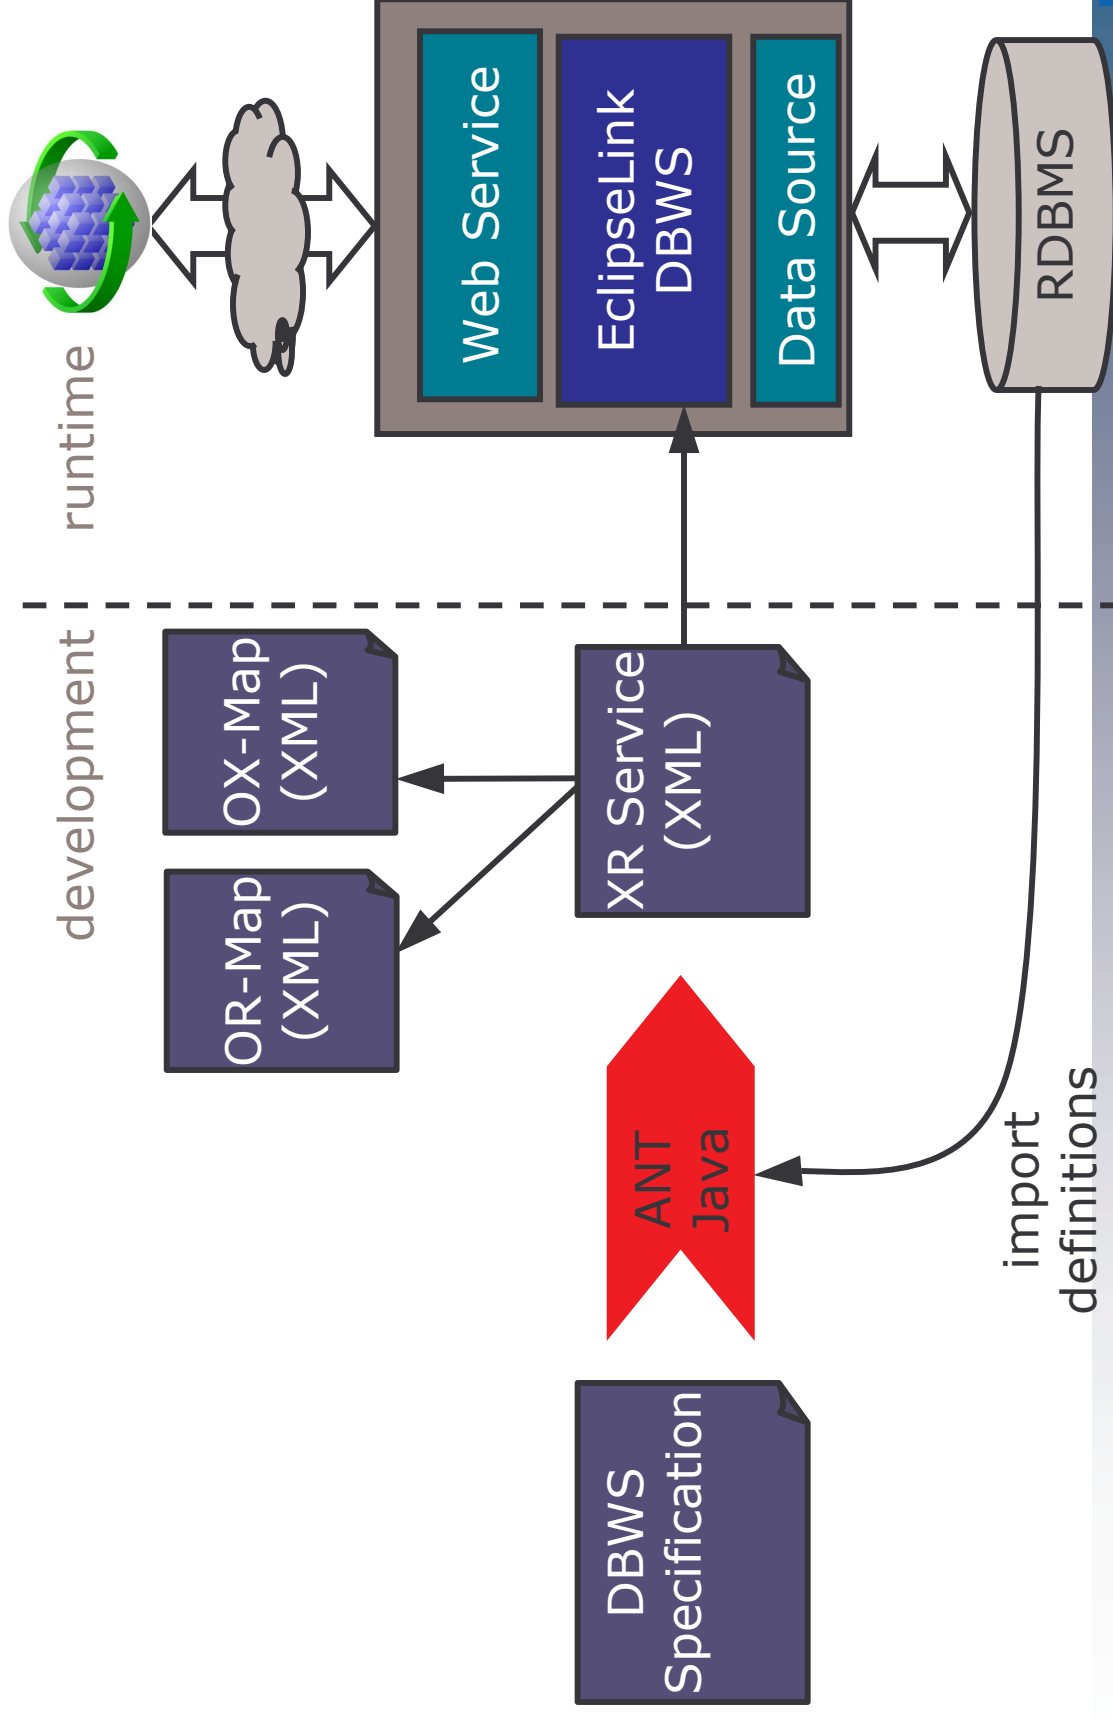
DEMO

EclipseLink MOXy (JAXB)

EclipseLink DBWS

- Simplified and efficient access to relational data through Web Services
- Minimal configuration with development utilities to retrieve metadata and generate/package Web Service
- Developers can fully customize the database access and XML mapping of the data
- Ideal for usage within SOA/SCA

EclipseLink DBWS



EclipseLink SDO

- What can you do?
 - Marshall/Unmarshall objects to/from XML
 - Define Types/Properties programmatically or derive from XSD
 - Generate JavaBean classes from XSD
 - Advanced mapping support for greater flexibility
- Why would you use it?
 - Schema/Structure unknown at compile time
 - Declarative metadata based tools/frameworks
 - XML-centric applications, need open content support
 - Dynamic content user interfaces

EclipseLink EIS

- Provide persistence support for non-relational data stores using Java EE Connector Architecture (JCA)
- Mapping interaction inputs and outputs to persistent domain model
 - XML mapping leveraging EclipseLink MOXy
 - Common Client Interface (CCI) mapping
- Visual mapping Workbench support
- Out of the box support for:
 - MQSeries, OracleAQ, Sun JCA, XML Files


EclipseLink and OSGi

- Work with OSGi expert group to define OSGi persistence services blueprint
- Deliver EclipseLink as OSGi bundle(s)
- Show through examples how to leverage within an OSGi solution
- Address technical challenges as a community

Combining Services

- Metadata based approach allows the same domain model to be mapped with multiple persistence services
 - Supports usage within Web Services/SOA/SCA
 - Domain model can be shared between persistence services (JPA, MOXy, EIS)
 - Transformations are bidirectional:
 - Unmarshall XML to objects and then persist
 - Marshall persistent objects to XML

Common Domain Model



EclipseLink and Spring

- EclipseLink JPA
 - Container
 - Template
- EclipseLink Native ORM Template
- EclipseLink MOXy
 - Direct, Spring WS, Spring Remoting, ...
- and many more possibilities...
 - Spring Batch, Spring OSGi, ...

EclipseLink JPA in Spring

```
@Repository
@Transactional
public class EntityManagerClinic implements Clinic {

    @PersistenceContext
    private EntityManager em;

    public Collection<Owner> findOwners(String lastName)
        throws DataAccessException
    {
        Query query =
            em.createNamedQuery("Employee.findOwners");
        query.setParameter("lastName", lastName + "%");
        return query.getResultList();
    }
}
```

EclipseLink in the Eclipse Ecosystem

- Provide an Eclipse persistence solution easily consumable by any project
 - Storage of metadata in RDBMS, XML, EIS
 - XML Messaging infrastructure
- Eclipse Projects
 - Dali JPA Tooling Project
 - Teneo to use EclipseLink for EMF model persistence
 - Maya for storage of deployment configuration
 - SOA Project for EclipseLink SDO

Where are we going?

- Delivery of initial 0.1-incubation milestone
 - Build and testing processes
 - Initial contribution functional
 - Spring Framework support
- Specifications: JAXB 2.0, SDO 2.1
- OSGi packaging and usage examples
- Database Web Service (DBWS)
- Data Access Service (DAS) - SDO with JPA
- Simplified DataMap Access and Dynamic Persistence

How can you get involved?

- Users
 - The 0.1-incubation milestone will be available soon
 - Try it out and provide feedback
 - File bug reports and feature requests
- Contributors
 - Contribute to roadmap discussions
 - Bug fixes
- Committers
 - Very interested in growing committer base

EclipseLink Summary

- First comprehensive Open Source Persistence solution
 - EclipseLink JPA: Object-Relational
 - EclipseLink MOXy: Object-XML
 - EclipseLink SDO: Service Data Objects
 - EclipseLink DBWS: Database Web Services
 - EclipseLink EIS: Non-Relational using JCA
- Mature and full featured
- Get involved

More Information

- www.eclipse.org/eclipselink
- Newsgroup: eclipse.technology.eclipselink
- Wiki: wiki.eclipse.org/EclipseLink
- Mailing Lists:
 - eclipselink-dev@eclipse.org
 - eclipselink-users@eclipse.org
- Blogs
 - Committer Team blog: eclipselink.blogspot.com
 - Doug's blog: java-persistence.blogspot.com

